

Pentaho Update

Bob Kemper
EVP Engineering, Pentaho

- Year in review
- Pentaho Strategy
- Hitachi and Pentaho
- Roadmap(s) - abbreviated

- Company awareness - Pentaho widely recognized as leader in Big Data
 - Number 1 or 2 in most “Share of Voice” metrics
- Major new large customer wins
 - Enterprise deployments and Big Data represent significant percent of business
- Continued staff growth in both US and EMEA teams
 - Across all areas of the business – Sales, marketing, services, support, engineering

- Major product releases
 - Several large product release during the year
 - 5.3 in Q1, 5.4 in Q2
 - 6.0 Platform release in October
- And ... Acquired by Hitachi Data Systems in May 2015
 - More on that to come....

Product Releases – 5.3 (February)

CE/EE Features

- Support for CDH 5.2 and MapR 4.0.1
- Named Clusters
- Carte Clustering
- PDI SDK as an “Official” release package
- BA Server / Plugin Improvements (Data Access for example - 0 bad builds during ALL 5.3.0 RC builds)
- PDI SDK documentation
- Carte Clustering documentation
- MT Documentation
- Over 300 Closed and Validated jiras! (Majority common code)

EE Only Features

- Using an existing model for publish
- Using user-defined annotations to assist auto-model generation
- Validated to work with Amazon Redshift and Impala
- Vizor as a package
- Analyzer JS API and Documentation
- PIR Improvements
- Analyzer API documentation

CE/EE Features

- New PDI Look and Feel (CE)
- BA Server Utilities
- C*Tools support for RequireJS
- Improved Named Clusters support
- AWS EMR Shim Support
- Spark Integration
- User Console localizations
- Mongo 3.0 Support
- Google Analytics step upgrades

EE Only Features

- Analyzer Extensibility APIs
- DI Server REST APIs
- SAP Hana support in PDI
- SDR Modeling improvements

Product Releases – 6.0 (October)

CE/EE Features

- PDI Data Services – Opened and major advances
- OSGI Support
- New PDI Marketplace
- Testing framework for Named Clusters
- Library and back end upgrades (Java8, Tomcat8, Drools, VFS, etc, etc.)
- REST Apis for User Management
- Groundwork for a “Merged BA and DI Server”
- Major step improvements/defect resolution in PDI
- Lineage capture in PDI (see blogpost by Pedro on displaying through yEd)

EE Only Features

- Inline Modeling in Analyzer
- Annotation improvements in PDI
- System monitoring via SNMP
- Lineage information through MetaIntegration
- Improved import/export

Pentaho Strategy

Modern, Cohesive Business Analytics and Data Integration Platform

- Full spectrum of analytics for all key roles
- Embeddable, cloud-ready analytics
- Broadest and deepest big data integration

Innovation Through an Open Heritage

- Open, pluggable, purpose-built for the future
- Sustained leadership in big data ecosystem

Business Momentum

- Over 1,500 commercial customers
- Over 10,000 production deployments



Pentaho is Architected for the Future

Purpose-built to support an evolving data & analytics landscape

Data Integration

Data Orchestration

Big Data

Data
Warehouse

Other

Business Analytics

Embedded Analytics

Analyzer

Dashboards

Reporting

Open



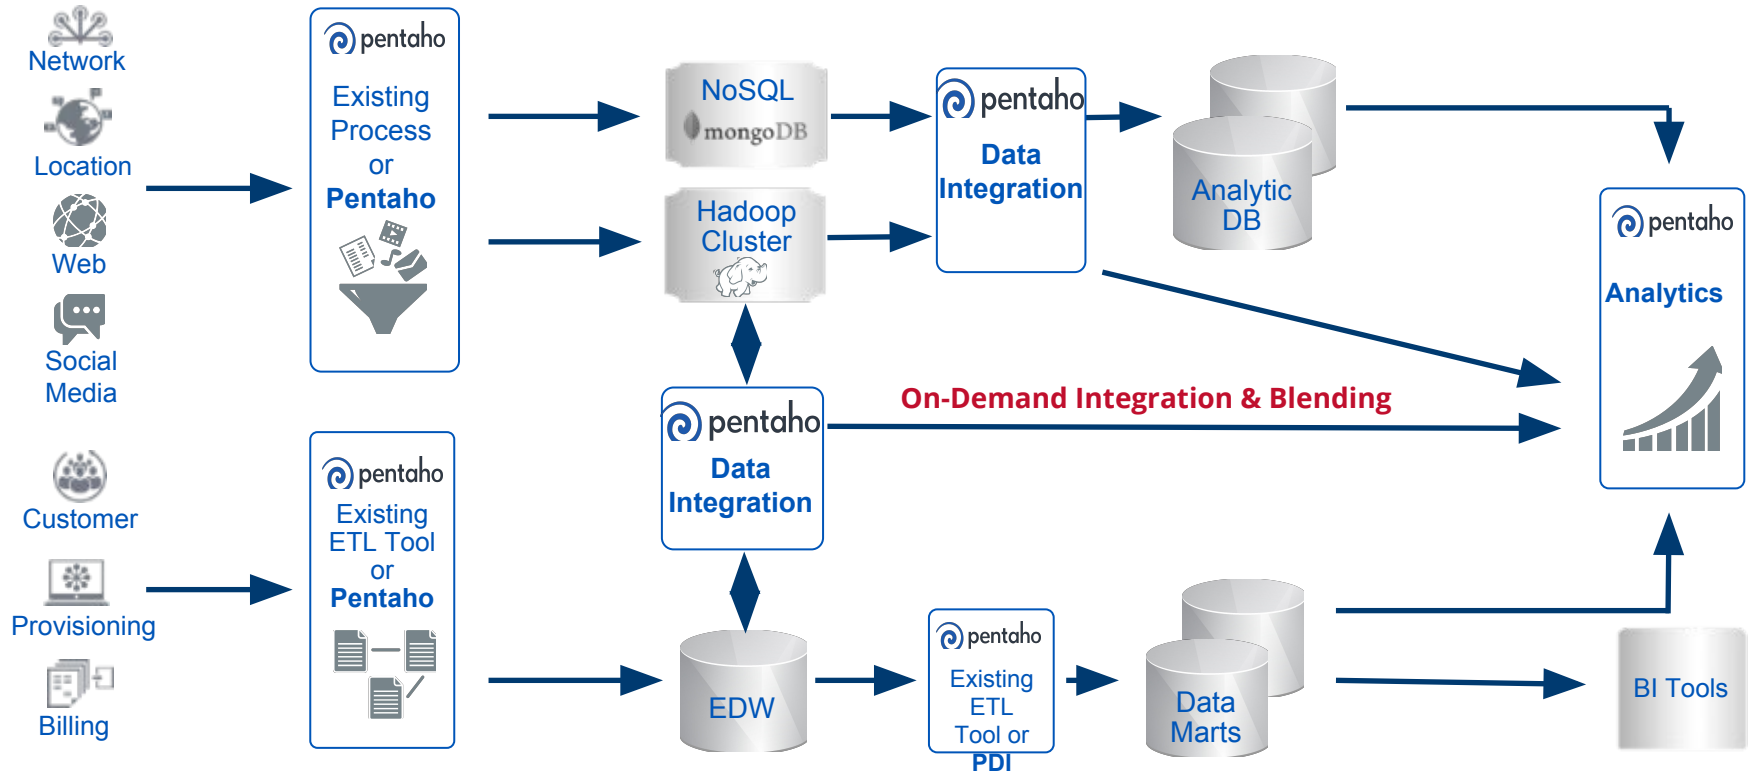
Pluggable



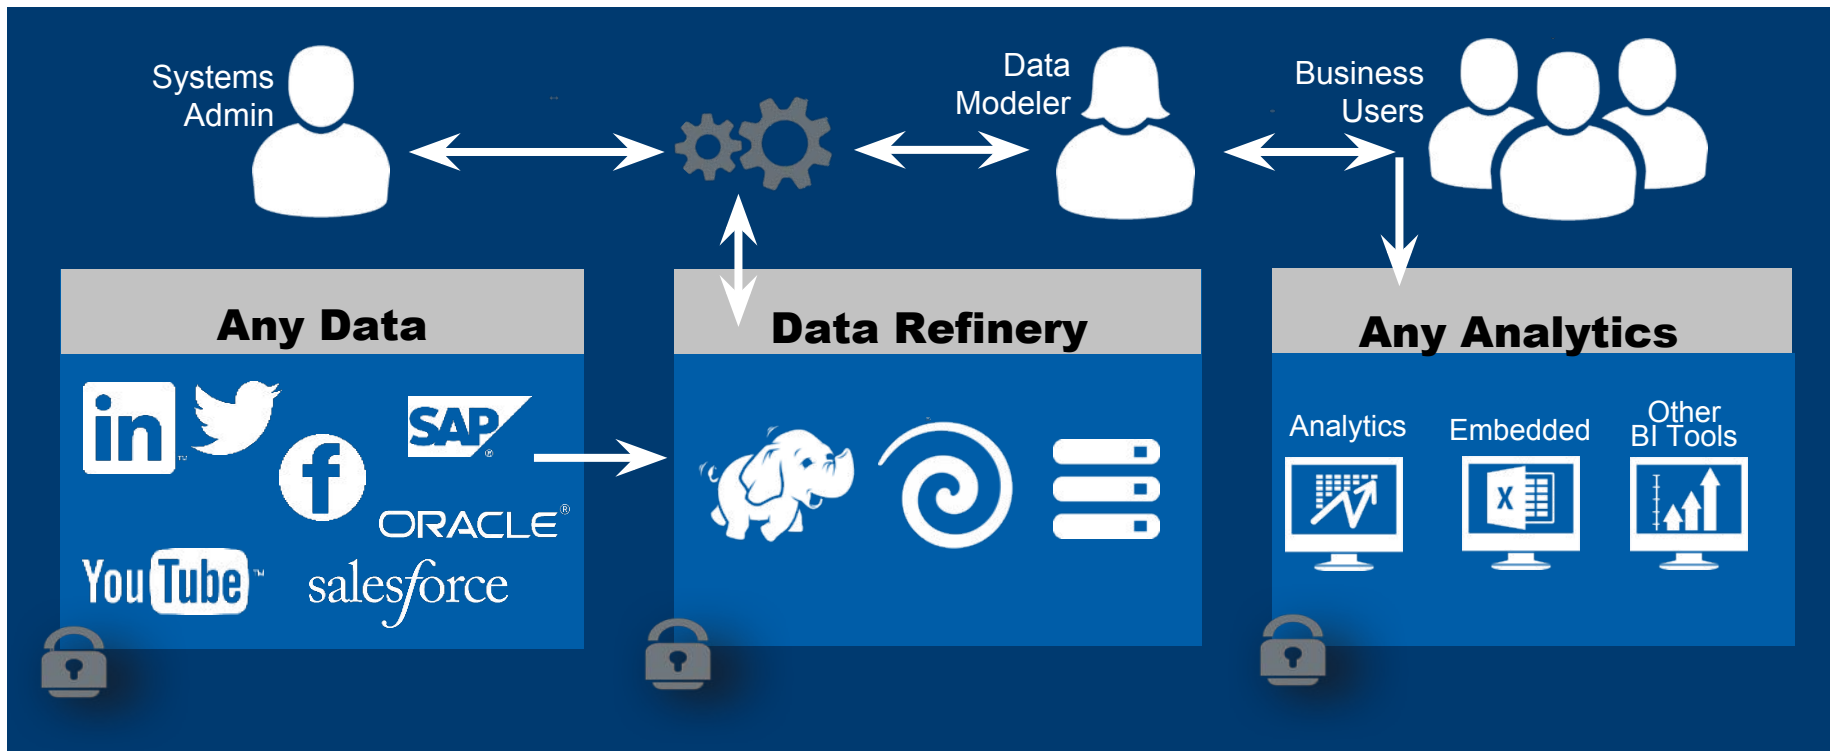
100% Java

Bringing it All Together

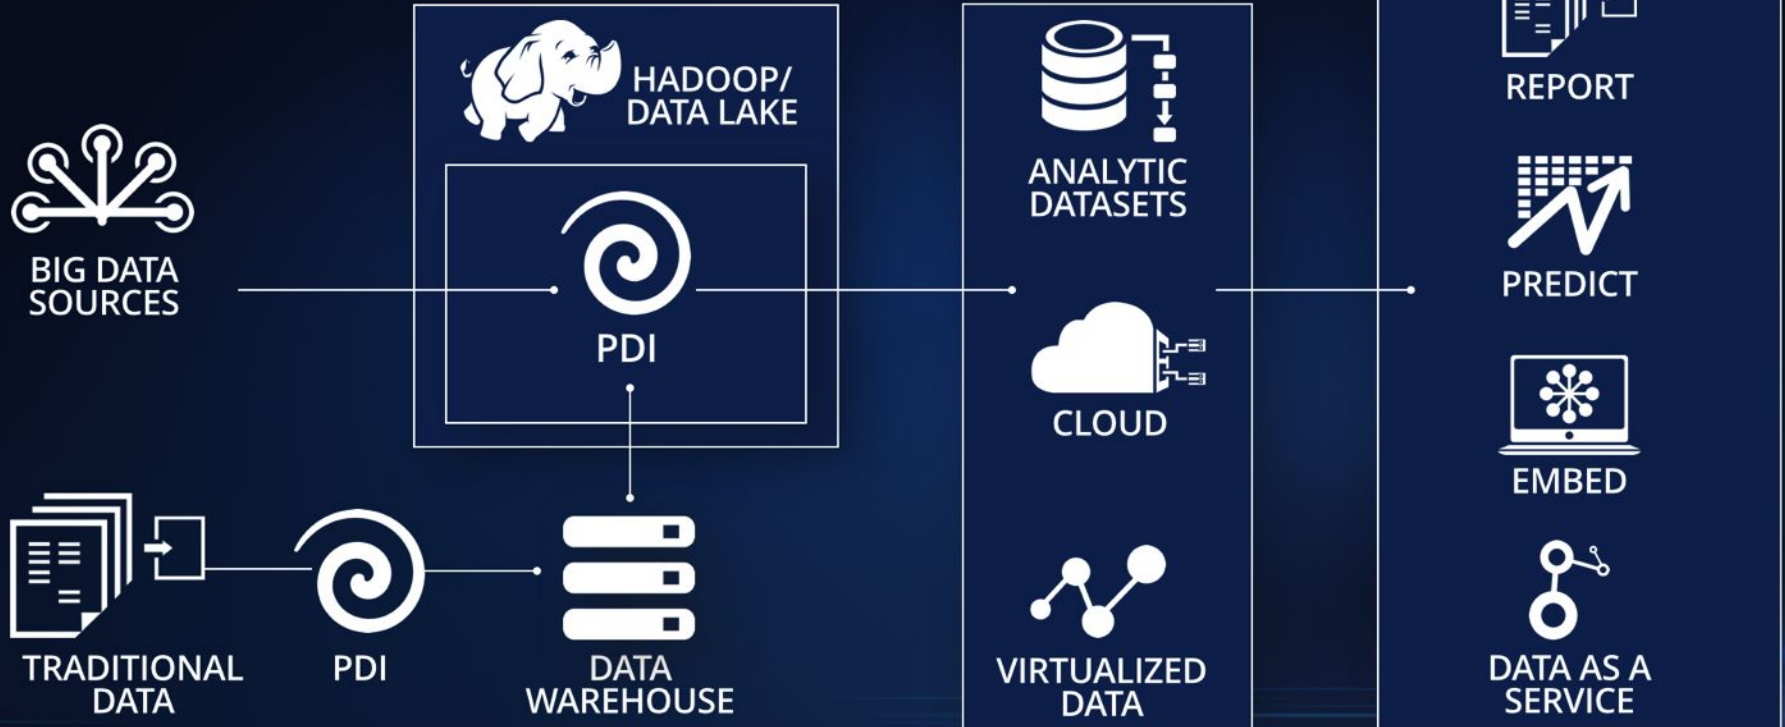
Pentaho Bridges Big Data and Traditional Architectures



End-to-End Self-Service Delivery of Big Data Analytics



BIG DATA LAKE: STRATEGIC CORPORATE ASSET



Why Did Hitachi Acquire Pentaho?



Leader in the **big data & business analytics** market

The Next Revolution


INDUSTRIAL



DIGITAL



INTERNET



INTERNET OF THINGS



Internet of Things in 2020



152M
connected
cars
(IHS
Automotive)



1.1B
smart
meters
installed
(Navigant
Research)



200B
connected
objects
(Intel)



\$19T
economic
value
created by
Internet of
Things
(Cisco)



\$2.5T
annual
impact
of IoT in
healthcare
(by 2025,
McKinsey)



>90%
IoT data to
be hosted
on cloud
platforms
(IDC)

The Opportunity in 2020



Internet of Things

\$7.1 trillion

Social Innovation

\$2 trillion

Integrating IT and Infrastructure Technologies for a Safer, Smarter, Healthier World



HITACHI'S SOCIAL INNOVATION BUSINESS

The Internet of Things *That Matter*

The Internet of Things *That Matter*

ENERGY



**MICROGRIDS &
GRID STABILITY**



**WIND
TURBINES**



**INFORMATION &
CONTROL SYSTEMS**



**ENERGY SAVINGS
AS A SERVICE**



**SOLAR
POWER**



**NUCLEAR POWER
PLANTS**



**BATTERY
SYSTEMS**

HEALTHCARE



**CLINICAL DATA
MANAGEMENT**



**NURSING CARE
SOLUTIONS**



**PARTICLE BEAM
THERAPY**



**PRECISION
MEDICINE**



**PHARMACEUTICAL
SOLUTIONS**



BIOINFORMATICS



**TESTING AND
DIAGNOSTIC SYSTEMS**



**MEDICAL IMAGING AND
DISTRIBUTION SYSTEMS**

PLATFORM AND TECHNOLOGY

Pentaho
Common analytics framework
Converged Platforms
Content Platform

HITACHI PATENTS

A leader in data analytics
technology patents

2800+ pending

PEOPLE AND EXPERIENCE

\$1B

IoT and big data R&D

500+ data scientists and
researchers worldwide

600+ employees in the
Social Innovation Business Unit

75+ resources in Hitachi Live Insight
Center of Excellence



Robotics



Traffic Analysis



Image Search



Power Grid Optimization



Human Behavior Analysis



Oil Field Analysis




Object Recognition




Mining Optimization



Facial Recognition



Augmented Reality



People Flow Analysis



Delivering Real IoT Solutions



SMART CITY



TELECOM



HEALTHCARE



**I.T.
ANALYTICS**



ENERGY



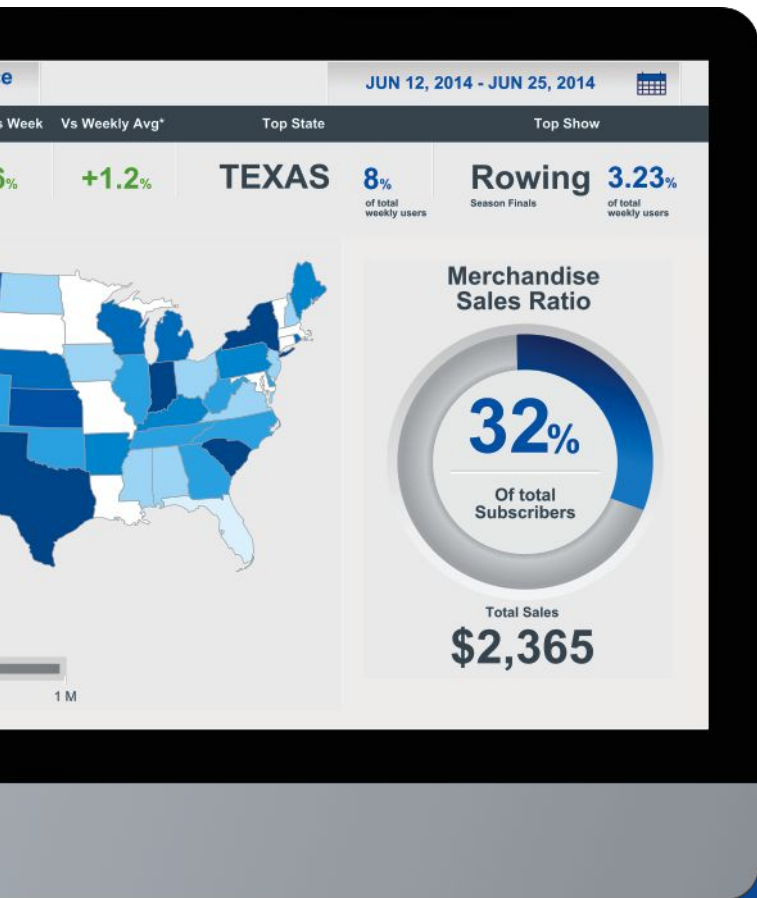
**INTELLIGENT
VEHICLE**



PENTAHO



IoT ANALYTICS PLATFORM



Pentaho Roadmap

(snapshot)

- The forward-looking statements contained in this document represent an outline of our current intended product direction. It is provided for information purposes only and is not a commitment to deliver any new or enhanced product or functionality, or that we will pursue the product direction described. Facts and circumstances may occur which may impact current plans, resulting in changes to the information in this presentation. This information is current only as of the date it is made and should not be relied upon in making purchasing decisions. The development, release (if at all), and timing of any features or functionality described for the Pentaho products remains at the sole discretion of Pentaho.

The challenge:

- More users have access and need access to data blending
- Security, governance and audit are recognized as critical for ALL Data
- Data from the lake needs to be blended with the EDW and external sources
- You can't always know what the data will look like until it arrives

The direction:

- Continue work on security & governance integration
- Improve ease of use and performance of data blending
- Bring data blending closer to the user
- Extend existing capabilities for dynamic blending

Continue security, governance & operations management

2015

- Kerberos support
- Lineage
- SNMP monitoring
- OPS Mart improvements

2016

- Flexibility in delegation of authority and action via Kerberos
- Focus on audit, logging, debugging, resource management

Futures

- Interoperability with security & governance providers

Improve Ease of Use and Performance of Data Blending

2015

- Sqoop step updates
- Spark executor step
- ElasticSearch updates

2016

- JSON and HBASE step updates for MapReduce
- Investigating broader support for emerging SQL on Hadoop technologies – Presto, Phoenix, Spark-SQL...

Futures

- Constant investigation into opportunities for performance and scaling
- Continue labs investigation of emerging Big Data technologies and productize as appropriate

Bring Powerful Big Data Blending Closer to the Analyst

2015

- Annotations stream step
- Auto-model build step
- Data Services

2016

- Easy to create prompts for dynamic transformations
- Investigations into faceted search workflows
- Best practices and samples for SDR use cases

Futures

- Investigation into web-based interface for blending & analytics together

Extend existing capabilities for dynamic & model driven transformations

2015

- Dynamic Blending (Metadata injection) exists and supports a variety of IOT, dynamic and model driven uses cases

2016

- Extend support & documentation of metadata injection
- Investigate IOT device and machine log formats to add

Futures

- Continue documentation and best practices for dynamic transformation with metadata injection
- Investigate input steps to add in support of IOT use case

The Challenge:

- Deployment automation & configuration management are increasingly popular – but there are many technologies at play
- Security, governance and audit are increasing concerns
- Use of analytics is exploding & increasing demand for easier application maintenance

The Direction:

- Improve deployment, configuration & upgrade
- Continue expanding capabilities for governance & operational management
- Standardize & simplify common role based administration tasks from a central location

Improve deployment, configuration & upgrade

2015

- Simplification of configuration with Hadoop & Data Services

2016

- Platform deployment, configuration & upgrades on local & remote servers
- Lifecycle management/versioning

Futures

- Investigating tenant administration & configuration

Continue expanding capabilities for governance & operational management

2015

- SNMP Monitoring
- Data Lineage in PDI
- OPS Mart improvements

2016

- Investigating next steps on lineage
- Focus on audit, logging, debugging, resource management

Futures

- Investigating specific industry regulatory requirements

Standardize & simplify common role based administration tasks from a central location

2015

- REST APIs for automated user management

2016

- **Phase 1:** Standardize & simplify common role-based administration tasks
- Investigate scheduling interoperability improvements

Futures

- **Phase 2:** Centralize administration & management
- Investigating improvements to tenant administration & management

The Challenge:

- The expectation of agile, tailored experiences is increasing
- Security, governance and audit are increasing concerns
- Data from the lake needs to be blended with the warehouse and other external sources

The Direction:

- Continue to expand and harden Analyzer APIs
- Invest in built-in capabilities
- Bring Big Data Blending closer to the user
- Extend and harden the Pentaho Platform Infrastructure

Expanding and hardening APIs

2015

- REST API for User Management

2016

- Visualization APIs
- Geo-mapping APIs

Futures

- Investigating componentization for lighter web embedding

Built-in Functionality

2015

- Inline modeling v1

2016

- Updates to visualizations
- Inline modeling for user defined groups and hierarchies
- Export formatting

Futures

- Investigating technologies for higher performance analytics

Bring powerful big data blending closer to the analyst

2015

- Annotations stream step
- Auto-model build step
- Data Services

2016

- Easy to create prompts for dynamic transformations
- Investigations into faceted search workflows

Futures

- Investigation into web-based interface for blending & analytics together

Extend and harden the Pentaho Platform Infrastructure

2015

- Spring 4 enables SSO and Third Party Authentication
- Java, Tomcat and JCR updates for security & performance

2016

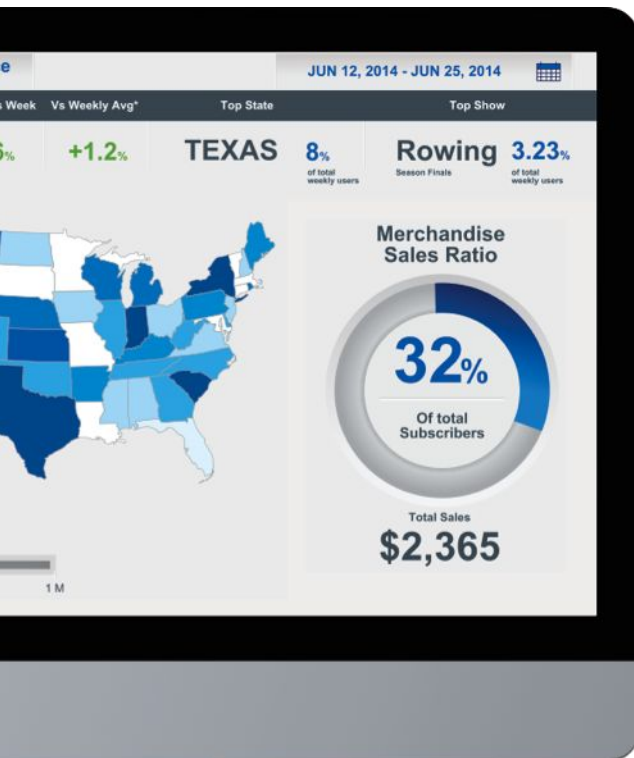
- Lineage available to enterprise metadata management tools
- Phase 1: simplified administration & management

Futures

- Investigating improvements to tenant administration & management

- Will be provided by Pedro –
 - STAY TUNED!

- Near term - Open the Annotations and Model build/publish steps
- New processes for community contributions
 - See Marc Batchelor's session later today
- High level agreement to open PIR
- Discussions with Hitachi and HDS about opening some of their tech
 - Candidates – streaming engine, search and content indexing, IoT platform components, etc.



- **A future-proof data foundation**
- **Innovation in the big data ecosystem**
- **Enterprise class platform**
- **Become MORE open – with full Hitachi support**
- **Embrace our roots of Open Source**

**THANK
YOU**